

In-Depth DNS Flow: "www.google.com" Request

Step-by-Step Breakdown

1. User enters www.google.com in a browser

- The browser checks its internal DNS cache.
- If not found, it proceeds to the next level.

2. Operating System DNS Cache

- The OS (like Windows or Linux) maintains its own cache.
- If a cached entry is found and still valid (TTL not expired), it's used.
- If not, the OS sends a DNS query to the configured **DNS Resolver** (typically from ISP or a public DNS like 8.8.8.8).

3. DNS Resolver (Recursive Resolver)

- Receives the query: "What is the IP of www.google.com?"
 - Checks its own cache.
 - If not found, begins the recursive resolution process by querying the DNS hierarchy:
-

Recursive Resolution Journey

4. Query to Root DNS Server

- Resolver contacts a **Root Server** (there are 13 sets globally: A to M).
- Root server doesn't know the IP but replies with the address of the responsible **TLD (Top-Level Domain) server**, e.g., .com.

5. Query to TLD Server

- Resolver asks the .com TLD server: "What's the authoritative DNS for example.com?"
- TLD replies with the IP of the **Authoritative Name Server** for example.com.

6. Query to Authoritative Name Server

- Resolver asks: "What is the IP address of www.google.com?"
 - Authoritative server responds with the actual IP address (e.g., 93.184.216.34).
 - This may include A (IPv4) or AAAA (IPv6) records, along with TTL.
-

Back to the Client

7. Recursive Resolver stores the response

- Caches the IP and TTL.
- Sends the IP back to the user's OS.

8. OS stores it in cache and gives it to the browser

9. Browser initiates TCP/HTTPS connection

- Now the browser connects to 93.184.216.34 over port 443 (HTTPS) or port 80 (HTTP), and fetches the web page.